

AI LITERACY

GLOSSARY

A plain-language guide to the terms people encounter most often when reading about or using artificial intelligence today, listed alphabetically.

A note on fast-moving terms

AI terminology moves quickly. Most terms below describe stable, well-established concepts, but a few (marked †) cover newer or still-evolving parts of the field — worth a quick check against current sources before you rely on them in a fast-moving context.

PLAIN LANGUAGE / BUSINESS USE / QUICK REFERENCE

A practical reference from ilene.ai

AI Literacy Glossary

A B C D E F G H I J K L M N O P Q R S T U V W X Y Z

A

Agent / Agentic AI

AI systems that can take multi-step actions toward a goal with limited supervision — browsing the web, using software tools, writing and running code — rather than just answering a single question.

AGI (Artificial General Intelligence)

A hypothetical AI with human-level ability across virtually all intellectual tasks, not just narrow specialties. No AI system has achieved this, and experts disagree on when — or whether — it will happen.

AI Ethics / Responsible AI

The study and practice of developing AI in ways that are fair, transparent, accountable, and beneficial — covering issues like bias, privacy, job displacement, and environmental impact.

AI Governance

The policies, processes, and oversight structures — inside a company or at a national or international level — that guide how AI is developed, deployed, and monitored. Inside a business, this often means approval workflows for new AI tools, usage policies, and audit trails.

AI Safety

The broad field concerned with preventing AI systems from causing harm as they become more capable and autonomous. It includes alignment but also covers misuse, accidents, and unintended behavior more generally.

Algorithm

A set of step-by-step instructions a computer follows to solve a problem. In everyday usage (“the algorithm showed me this video”), it usually refers to the recommendation systems used by social media and streaming platforms.

Alignment

The challenge of making AI systems reliably do what humans intend and value — helpful, honest, and harmless — rather than pursuing goals in unintended or harmful ways.

API (Application Programming Interface)

A defined way for one piece of software to talk to another. In AI, “using the API” usually means a developer is connecting their own product to a model like Claude or GPT programmatically, rather than through a chat window.

Artificial Intelligence (AI)

A broad field of computer science focused on building systems that can perform tasks normally requiring human intelligence, such as understanding language, recognizing images, or making decisions. “AI” is an umbrella term — most of what people call AI today is a specific type called machine learning.

Attention Mechanism

A technique that lets a model weigh how relevant different parts of an input are to each other when producing an output. It's the core idea behind the Transformer architecture that underlies modern language models.

B

Bias

Systematic unfairness in AI outputs, usually inherited from patterns in training data. For example, a hiring tool trained on past decisions may replicate historical discrimination.

Black Box

A description of the fact that even AI developers often cannot fully explain why a model produced a specific output. The field working on this problem is called interpretability, or explainable AI (XAI).

C**Chain-of-Thought Reasoning**

A technique where a model works through intermediate reasoning steps — shown or hidden — before giving a final answer, which often improves performance on complex math, logic, or multi-step problems.

Chatbot / AI Assistant

A conversational interface to an AI model. The chatbot is the product; the model is the engine underneath.

Compute †

Shorthand for the computing power (usually specialized chips called GPUs) needed to train and run AI. Access to compute is a major factor in who can build advanced AI, and demand for it shifts quickly as the field evolves.

Constitutional AI

A training approach, developed by Anthropic, where a model is guided by a written set of principles rather than relying only on humans rating individual outputs one by one — intended to make alignment more scalable and transparent.

Context Window

The amount of text (measured in tokens) a model can consider at one time — including your conversation history and any documents you share. If a conversation exceeds the context window, the model loses track of earlier parts.

Copilot

A general term for AI tools that assist a person inside an existing workflow rather than replacing it — most commonly associated with AI-assisted coding, though the term is also used more loosely (and as a specific product name by some vendors) for AI features embedded in everyday software.

D**Data Labeling / Annotation**

The process of tagging raw data — images, text, audio — with the correct answers or categories so it can be used to train a supervised machine learning model. Often done by large teams of human reviewers.

Data Privacy

The question of what happens to information you share with an AI system — whether it is stored, reviewed by humans, or used to train future models. Policies vary by company and product; check before sharing sensitive information.

Deep Learning

A subset of machine learning that uses neural networks with many layers. It powers most modern AI breakthroughs, including image recognition, speech recognition, and language models.

Deepfake

Synthetic media — usually video, images, or audio — that realistically depicts a real person saying or doing something they never did. A major concern for misinformation and fraud.

Diffusion Model

A generative AI technique, most often used for images, that learns to create content by starting from random noise and gradually refining it into a coherent output over many steps.

E

Embedding

A way of representing words, sentences, images, or other data as a list of numbers (a vector) that captures meaning, so that similar concepts end up numerically close together. This underpins search, recommendations, and retrieval-augmented generation.

Emergent Abilities †

Capabilities that show up in large models but were absent or much weaker in smaller versions of the same model — appearing as a byproduct of scale rather than being deliberately engineered. Whether this is a sharp, sudden effect or a smoother one that just looks sudden on certain metrics is still actively debated among researchers.

F

Fine-Tuning

Taking a model that has already been trained and giving it additional, more specialized training — for example, adapting a general language model to answer medical questions or follow instructions more reliably.

Foundation Model

A large, general-purpose model trained on broad data that can be adapted for many different applications — the “foundation” other products are built on.

G

Generative AI

AI that creates new content — text, images, audio, video, or code — rather than just classifying or analyzing existing content. Chatbots and image generators are common examples.

GPU (Graphics Processing Unit)

A chip originally designed for rendering video game graphics that turned out to be well-suited to the mathematics of training neural networks. Demand for GPUs has made chipmakers central players in the AI industry.

Grounding

The practice of tying an AI's output to specific, verifiable sources — documents, search results, a database — rather than letting it rely solely on patterns learned during training. A well-grounded system still needs its sources checked, but it reduces hallucination.

Guardrails

Safety measures built into AI systems to prevent harmful outputs, such as refusing to give dangerous instructions or generate abusive content.

H

Hallucination

When an AI confidently generates false information — invented facts, fake citations, nonexistent quotes. This happens because language models predict plausible-sounding text rather than looking facts up in a database. Always verify important claims from AI.

Human-in-the-Loop

An approach where a person reviews, approves, or corrects an AI system's output before it takes effect, rather than letting the AI act fully autonomously. Common in high-stakes uses like hiring, healthcare, and finance.

I

Inference

The act of running a trained model to get an output — every time you ask a chatbot a question, it performs inference. Distinct from training, which happens before the model is deployed.

J

Jailbreak

An attempt to trick an AI into bypassing its own safety rules through cleverly crafted prompts, aimed directly at the model itself.

K

Knowledge Cutoff

The point in time after which a model has no training data. A model doesn't know about events after its cutoff unless it can search the web, which is why AI assistants can give outdated answers on current topics.

L

Large Language Model (LLM)

A type of generative AI trained on enormous amounts of text to predict and generate language. LLMs power chatbots and assistants. “Large” refers to both the amount of training data and the number of parameters in the model.

Latency

The time it takes an AI system to respond after receiving a request. Lower latency matters more for real-time uses, like voice assistants or live customer support, than for tasks where a short wait is acceptable.

M

Machine Learning (ML)

A method of building AI where a system learns patterns from data rather than following rules written by a programmer. Instead of being told “a cat has pointy ears and whiskers,” the system is shown many cat photos and learns to recognize cats on its own.

Misinformation / Synthetic Content

AI makes it cheap and fast to produce convincing false text, images, and audio at scale. Media literacy now includes asking whether content might be AI-generated.

Mixture of Experts (MoE)

A model architecture made up of many smaller specialized subnetworks (“experts”), where only a subset are activated for any given input. This lets very large models run more efficiently than if every part were used every time.

Model

The end product of training — the trained system itself. When people say “the model,” they mean the AI program that takes inputs (like your question) and produces outputs (like an answer).

Model Context Protocol (MCP) †

An open standard for connecting AI assistants to external tools, files, and data sources in a consistent way, rather than every app building its own custom integration. This is a genuinely new and still-evolving part of the AI ecosystem, so specifics are worth checking against current documentation.

Multimodal

Able to work with more than one type of input or output — for example, a model that can understand both text and images, or generate audio as well as writing.

N

Narrow AI

AI designed for one specific task (translating languages, playing chess, detecting spam). All AI in practical use today is narrow AI, even very capable systems.

Natural Language Processing (NLP)

The branch of AI focused on enabling computers to understand, interpret, and generate human language. It's the broader field that today's language models are the most recent and powerful expression of.

Neural Network

A computing system loosely inspired by the human brain, made up of interconnected nodes ("neurons") organized in layers. Data passes through these layers, and the connections are adjusted during training until the network produces useful outputs.

O

Open Source / Open Weights

Models whose underlying files are released publicly so anyone can download, run, or modify them, as opposed to "closed" models accessible only through a company's service. "Open weights" is the more precise term when the training data and code aren't also released.

Overfitting

When a model learns the training data too specifically — including its noise and quirks — so it performs well on that data but poorly on new, unseen data. A core quality concept across all of machine learning, not just AI.

P

Parameters

The internal numerical settings of a model that get adjusted during training. Modern large models have many billions of parameters. More parameters generally means more capacity to learn, but not automatically better performance.

Prompt

The input you give an AI — your question, instruction, or request. The quality and clarity of a prompt strongly affects the quality of the response.

Prompt Engineering

The practice of crafting prompts to get better results from AI systems, such as giving examples, specifying a format, or breaking a task into steps.

Prompt Injection

An attack where hidden or disguised instructions are embedded in content an AI processes — a webpage, document, or email — to manipulate it into ignoring its original instructions or taking unintended actions. Different from jailbreaking, which targets the AI directly rather than through inserted content.

R

RAG (Retrieval-Augmented Generation)

A technique where a model looks up relevant documents (from a database or the web) before answering, so it can respond with current or specialized information instead of relying only on its training data. This reduces, but does not eliminate, hallucination.

Reasoning Model †

A category of AI model designed to spend more computation working through a problem step by step before answering, often improving performance on complex math, coding, or logic tasks at the cost of speed. This is an actively developing area, and both the terminology and the techniques are still shifting.

Red Teaming

The practice of deliberately trying to find weaknesses, harmful outputs, or safety failures in an AI system, usually before release, so they can be fixed. Borrowed from cybersecurity and military terminology.

Reinforcement Learning from Human Feedback (RLHF)

A training technique where humans rate a model's outputs and the model learns to produce responses people prefer. It's a key method for making chatbots more helpful and less harmful.

S

Shadow AI

The use of AI tools inside an organization without the knowledge, approval, or oversight of IT or leadership — for example, an employee pasting sensitive company data into a public chatbot. A growing concern in enterprise AI governance, and a common finding in early-stage AI maturity assessments.

Small Language Model (SLM)

A language model built to be smaller and more efficient than large frontier models, often designed to run on a phone, laptop, or private server rather than in the cloud — trading some general capability for speed, cost, or privacy.

Synthetic Data

Artificially generated data used to train or test AI models, created either by another AI system or by simulation, rather than collected from real-world events. Used when real data is scarce, sensitive, or expensive to gather.

System Prompt

Hidden instructions given to a chatbot by its developer that shape its behavior, tone, and limits before you ever type anything.

T

Temperature

A setting that controls how random or predictable an AI's output is. Low temperature produces more focused, repeatable answers; high temperature produces more varied and creative, but less predictable, output.

Token

The unit of text a language model processes — roughly a word or piece of a word. “Understanding” might be split into “under” and “standing.” Models read and generate text token by token, and pricing and limits are often measured in tokens.

Training

The process of teaching a model by exposing it to large amounts of data and adjusting its internal settings until it performs well. Training large models can take months and enormous computing power.

Training Data

The examples a model learns from. For language models, this is typically huge amounts of text from books, websites, and other sources. The quality and diversity of training data strongly shape how a model behaves.

Transformer

The neural network architecture, introduced in 2017, that underlies almost all modern large language models. Its key innovation is the attention mechanism, which lets the model weigh the relevance of every part of the input to every other part at once, rather than processing text strictly in order.

Turing Test

A thought experiment proposed by Alan Turing in 1950: if a machine can converse so convincingly that a human cannot reliably tell it apart from a person, it passes. Still widely referenced, though most researchers no longer treat it as a meaningful measure of intelligence.

V

Vector Database

A specialized database built to store and quickly search embeddings — numerical representations of meaning — rather than exact text matches. A key piece of infrastructure behind RAG and AI-powered search.

W

Watermarking (AI Content) †

Techniques for embedding a signal, visible or invisible, into AI-generated text, images, audio, or video so it can later be identified as AI-generated. An emerging and still-developing tool for addressing misinformation, with no universal standard yet.

Z

Zero-Shot / Few-Shot Learning

A model's ability to perform a task with no examples (“zero-shot”) or just a handful of examples given in the prompt (“few-shot”), rather than needing to be specifically retrained for that task.

PUT IT INTO PRACTICE

Quick Habits for AI Literacy

01

Verify facts — treat AI output as a first draft, not a final source, especially for names, numbers, dates, and citations.

02

Know the cutoff — ask whether the AI's information could be outdated, and whether it has web search available to check.

03

Protect your data — avoid sharing sensitive personal or company information unless you understand the product's privacy policy — and know whether your organization has an AI usage policy at all.

04

Watch for synthetic media — a realistic image, voice, or video is no longer proof something happened.

05

Remember it predicts, it doesn't know — language models generate plausible text; confidence is not accuracy.

06

Ask who's accountable — for any AI-assisted decision that matters — a hire, a diagnosis, a financial call — there should be a clear human owner, not just “the AI said so.”

AI literacy is not about knowing every term. It is about knowing enough to ask better questions.

ABOUT THIS GUIDE

General information only. This guide is intended for AI literacy and does not constitute legal, regulatory, security, or professional advice. AI terminology, capabilities, vendor practices, and requirements change quickly; verify current source documentation and applicable requirements before making business decisions.